

Manyue Javvadi . Jarvis Consulting

I am a Databricks Certified Data Engineer with professional experience in software development and a strong foundation in building reliable, scalable data systems. I began my career as a Software Engineer at Cognizant, working on Java-based enterprise applications in the insurance domain, where I gained hands-on experience with structured development practices, SQL-based data access, and production support. With a bachelor's degree in Commerce and postgraduate training in Big Data and Applied AI, I bring both technical depth and a strong understanding of business processes and decision-making. I am particularly interested in designing data platforms that support analytics, reporting, and data-driven decisions in large organizations such as banks and enterprise environments.

Skills

Proficient: Databricks, Java, Python, SQL, PostgreSQL, Linux/Bash, Apache Spark, Git

Competent: Data Warehousing, ETL/Data Pipelines, Agile/Scrum, Docker, Apache Airflow

Familiar: Hadoop Ecosystem, Cloud Platforms (Azure/GCP), Machine Learning Concepts, Data Modeling, CI/CD Fundamentals

Jarvis Projects

Project source code: https://github.com/jarviscanada/jarvis_data_eng_ManyueJavvadi

Azure Databricks Data Engineering [GitHub]: Designed and delivered two production-grade data pipelines on Azure Databricks for a fintech client using Medallion Architecture (Bronze - Silver - Gold). The first pipeline ingests financial transaction data from Azure SQL Database via JDBC and ADLS Gen2 via External Location, applies cleaning and enrichment in a silver layer, and produces 14 fraud analysis aggregations in a gold layer powering a live AI/BI dashboard. The second pipeline uses Delta Live Tables to ingest daily stock data from Alpha Vantage API for four equity symbols, computes 7/30/90 day price and volume trend metrics, and refreshes a market intelligence dashboard on a daily schedule. Both pipelines are orchestrated using Databricks Jobs. Encountered and resolved real infrastructure blockers including vCPU quota exhaustion, ADF Unity Catalog compatibility limitations, and large nested JSON ingestion challenges.

Cluster Monitor [GitHub]: Built a Linux-based monitoring agent to collect CPU and memory usage data, persist system metrics in PostgreSQL, and enable historical analysis using SQL. Designed database schemas, ingestion scripts, and containerized the solution using Docker to support operational visibility.

Core Java Applications [GitHub]: Developed multiple Java backend applications including a grep-style command-line tool for efficiently searching large directory structures using regular expressions. Focused on clean interfaces, modular design, error handling, Maven-based builds, and containerized execution.

Python Data Analytics [GitHub]: Implemented an end-to-end analytics workflow using PostgreSQL and Python to analyze retail transaction data. Produced revenue trends, order and cancellation analysis, customer behavior insights, and RFM segmentation to support targeted marketing strategies.

Databricks PySpark Retail Analytics [GitHub]: Re-implemented the London Gift Shop retail analytics solution using Databricks and PySpark to handle larger-scale processing beyond a single-machine Pandas workflow. Ingested the retail dataset into Unity Catalog as a managed table, performed data validation (nulls profiling, cancellation detection, invalid quantity/price checks), and built key business insights using PySpark DataFrames and Spark SQL validation (invoice-level revenue, monthly KPIs, growth, and customer activity insights).

Highlighted Projects

Retail Customer Analytics Platform [GitHub]: Designed and delivered a customer analytics solution that transforms raw transactional data into actionable insights, including monthly sales trends, customer activity analysis, and RFM-based customer segmentation for marketing decision support.

DevAstroAI Full-Stack AI Platform [GitHub]: Built and deployed a production-grade AI-powered web application from scratch. Backend - FastAPI (Python) integrated with Swiss Ephemeris for complex astronomical computations, served via Railway. Frontend - Next.js deployed on Vercel. AI layer - Anthropic Claude API with structured multi-step prompt engineering to produce domain-specific bilingual (English + Telugu) analysis outputs. Features include real-time astronomical calculations, multi-topic inference engine, PDF report generation, and a multi-step date-finding algorithm.

Designed for extensibility with modular routers, knowledge base injection, and conversation history management. Validated by a domain expert for output accuracy.

Professional Experiences

Software Engineer, Cognizant Technology Solutions (Jun 2021 - Sep 2023): Developed and maintained Java-based enterprise applications in the insurance domain. Migrated multiple applications from SVN to GitHub and helped establish CI/CD workflows, improving development productivity. Wrote SQL queries and procedures for reporting and downstream integration, resolved production issues during system upgrades, and mentored new team members through documentation and knowledge sharing.

Data Engineer (Consultant), Jarvis Consulting Group (Nov 2025 - Present): Working on data engineering projects involving data ingestion, analytics, and pipeline development using Python, SQL, PostgreSQL, and containerized environments for enterprise use cases. Collaborating with senior engineers to deliver production-ready proof-of-concept solutions for clients.

Education

George Brown College (2024 - 2025), Postgraduate Certificate, Applied AI Solutions Development - Graduated with distinction (96%)

Georgian College (2024 - 2024), Postgraduate Certificate, Big Data Analytics - Graduated with high academic standing (86%)

SRM University (2018 - 2021), Bachelor of Commerce, Commerce - Graduated with strong academic performance (85%)

Miscellaneous

- Databricks Certified Data Engineer Associate (2025)
- AWS Academy Graduate Cloud Architecting (2024)
- IBM Enterprise Design Thinking Practitioner (2024)